

Computational
Social Science

Modeling Count Data

Roberto Cerina

05.03.2024

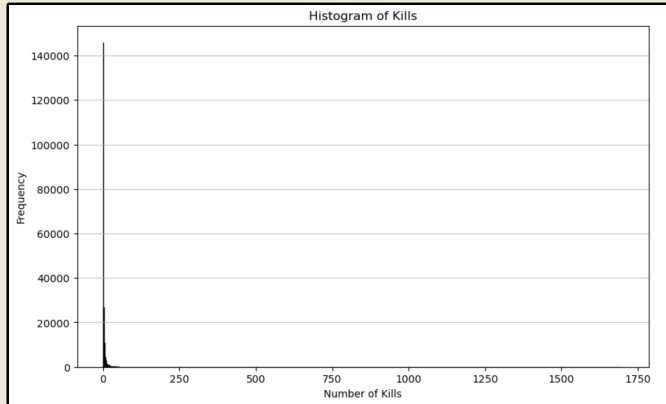


UNIVERSITY OF AMSTERDAM

Count Data

☞ $y_i \in \mathbb{N}_0$

☞ E.g [...26, 1, 1, 0, 1, 0, 0, 2, 1, 1, 1, 1, 1, 1, 0, 1, 0, 1, 1, 1, 1, 0, 0, 1, 3, ...]




Modeling Option 1: Log-Linear Regression

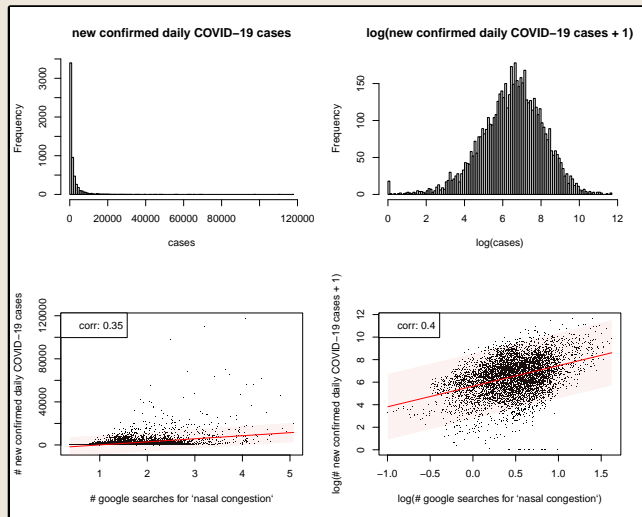
- Consider modeling counts as normally distributed...
 - ✖ *a.* predictions can be negative (not an issue if y is very large);
 - ✖ *b.* shape of predictive distribution might be poor due to non-normality (non-symmetry, non-constant variance) of errors;
 - ✖ *c.* predictions will not be integers;
- We can fix *a.* and *b.* by transforming the dependent variable onto the log-scale:
- Empirical posterior distribution of count-scale fitted values:
- ① Learn posterior distribution of β from model fit to $\log(y)$:
 $\log(y_i) \sim N(\beta_0 + \beta_1 x_{i1} + \dots + \beta_p x_{ip}, \sigma^2);$
- ② Simulate S values of β :
 $\beta^s \sim MN(\hat{\beta}, \hat{\Sigma});$
- ③ for each simulation, sample values of $\log(y_i)$:
 $\tilde{y}_i^s \sim N(\beta_0^s + \beta_1^s x_{i1} + \dots + \beta_p^s x_{ip}, \hat{\sigma}^2);$
- ④ Convert fitted values on the original scale:
 $y_i^s = \exp(\tilde{y}_i^s)$

Modeling Option 1: Log-Linear Regression

- Taking the log a random variable encourages 'normality' and 'constant variance'...
- \log_e stretches values from 0 to 1 into the negative, and *squishes* values above 1. The further away from 1 the larger the impact of the the log.



Modeling Option 1: Log-Linear Regression



Modeling Option 1: Log-Linear Regression

☞ Interpretation:

- **Linear model:**

$$y_i = \beta_0 + \beta_1 x_i + e_i$$

$$\frac{\partial y}{\partial x} = \beta_1$$

- A one-unit change in x is associated with a β_1 change in y .

Modeling Option 1: Log-Linear Regression

👁 Interpretation:

- **log(x) model:**

$$y_i = \beta_0 + \beta_1 \log(x_i) + e_i$$

$$\frac{\partial y}{\partial x} = \frac{1}{x} \beta_1$$

📖 $\log(x) + 1 = \log(x) + \log(e) = \log(ex) = \log(2.72x)$

- β is the expected change in y for a 172%¹ increase in x ;
- A general $p\%$ increase in x can be expressed as: $x \left(1 + \frac{p}{100}\right)$;

$$\begin{aligned}\Delta y &= \beta \log\left(x\left(1 + \frac{p}{100}\right)\right) - \beta \log(x) \\ &= \beta (\log(x) + \log(1 + \frac{p}{100})) - \beta \log(x) \\ &= \beta \log(1 + \frac{p}{100})\end{aligned}$$

$\therefore \beta \times \log(1 + \frac{p}{100})$ is the expected change in y for a $p\%$ increase in x .

- For small p , due to Taylor expansion², $\log(1 + \frac{p}{100}) \approx \frac{p}{100}$

\therefore if $p = 1$, $\frac{\beta}{100}$ is approximately the change in y for a 1% change in x .

¹For any factor c multiplied by x , the percentage change in x is $\Delta\% = 100 * (c - 1)$

² $\log(1 + x) = x - \frac{x^2}{2} + \frac{x^3}{3} - \frac{x^4}{4} + \dots$

Modeling Option 1: Log-Linear Regression

👁 Interpretation:

- **log linear model:**

$$\log(y_i) = \beta_0 + \beta_1 x_i + e_i$$

$$y_i = \exp(\beta_0 + \beta_1 x_i + e_i)$$

$$\frac{\partial y}{\partial x} = \exp(\beta_0 + \beta_1 x_i + e_i) \beta_1 = y_i \beta_1$$

📖 The rate of change in y_i for a one-unit change in x is proportional to the current value of y_i ;

🔗 A one-unit increase in x multiplies the outcome by $\exp(\beta_1)$:

$$\begin{aligned} \frac{y(x_i + 1)}{y(x_i)} &= \frac{\exp(\beta_0 + \beta_1(x_i + 1) + e_i)}{\exp(\beta_0 + \beta_1 x_i + e_i)} \\ &= \frac{\exp(\beta_1(x_i + 1))}{\exp(\beta_1 x_i)} = \\ &= \exp(\beta_1(x_i + 1) - \beta_1 x_i) = \\ &= \exp(\beta_1) \\ y(x_i + 1) &= y(x_i) \times \exp(\beta_1) \end{aligned}$$

- for any other size of change in x : $y(x_i + c) = y(x_i) \times \exp(c\beta_1)$

Modeling Option 1: Log-Linear Regression

☞ Interpretation:

- **log linear model:**

$$\begin{aligned}y(x + 1) &= \exp(\beta_0 + \beta_1(x_i + 1) + e_i) \\&= \exp(\beta_0 + \beta_1 x_i + e_i) \exp(\beta_1) \\&= y(x_i) \exp(\beta_1) \\&\approx y(x_i)(\beta_1 + 1)\end{aligned}$$

- \therefore a 1 unit increase in x is associated with a $\beta_1\%$ increase in y ;
- The approximation above is especially true when β is small;
- Check $\beta^2/2$ – this is the first-order error of the Taylor expansion, and should be ≈ 0 , for the approximation above to be reliable;
- As $\Delta x \rightarrow 0$, we recover the ‘instantaneous rate of change’ – or the derivative from the previous slide.

Modeling Option 1: Log-Linear Regression

👉 Interpretation:

- **log-log model:**

$$\log(y_i) = \beta_0 + \beta_1 \log(x_i) + e_i$$

$$\frac{\partial y}{\partial x} = \exp(\beta_0 + \beta_1 \log(x_i) + e_i) \frac{1}{x_i} \beta_1 = \frac{y_i}{x_i} \beta_1$$


- Looking at the instantaneous rate of change above, we can see the impact of a one-unit-change in x onto y can be interpreted as a mix of the partial models considered up to now:
- a $p\%$ increase in x is associated with an $\exp(\log(1 + \frac{p}{100}) \beta_1)$ percentage increase in y
- When β and p are small:
 $\approx \exp(\frac{p}{100} \beta_1) \approx \frac{p}{100} \beta_1 + 1$
- ✓ You get approximately a $p\beta_1\%$ increase in y for a $p\%$ increase in x .

Problems with *log-plus-1* Regression

- ➡ To analyse count data with a linear regression model we have to take $\log(y)$...
- ✗ ... $\log(0)$ is not defined ...
 - ... need to add a constant to y to shift it away from 0 ...
 - ... but this can introduce substantial bias to regression coefficients:


$$\begin{aligned}\log(y_i) &= \beta_0 + \beta_1 x_i + e_i \\ \log(y_i + 1) &= \lambda_0 + \lambda_1 x_i + e_i \\ \hat{\lambda}_1 &\approx \frac{\hat{y}(x = \bar{x})}{1 + \hat{y}(x = \bar{x})} \hat{\beta}_1\end{aligned}$$

Problems with *log-plus-1* Regression



Modeling Option 2: Poisson Regression

- ⇒ Log-linear regression uses a mathematical expedient to ‘massage’ the data into a convenient form, to analyse with methods we are familiar with...
- ⇒ arguably the log-linear model does not represent the ‘true’ data generating process –
- ✗ counts from skewed samples are typically not ‘really’ normally distributed....
- ⇒ The default distribution for count data is the Poisson distribution !



The 'Homogeneous Rate' Poisson Model

$$y_i \sim \text{Poisson}(\lambda) \quad \forall y_i \in \mathbb{N}_0, i \in \{1, \dots, n\}$$

⇒ $\lambda \in (0, \infty)$ – *rate* parameter

- expected number of events per 'interval' (e.g. per time- or space-units)
- the Poisson model assumes the variance of the data is equal to the expected value:

$$E[y] = \text{Var}[y] = \lambda;$$

- this is different from the normal (and log-normal) models, which allow for estimation of separate mean and variance...

$$E[y] = \mu$$

$$\text{Var}[y] = \sigma^2$$

The 'Homogeneous Rate' Poisson Model



Note:

- 'square' nodes indicate 'observed / known values';
- 'circular' nodes indicate 'unknown parameters';
- 'solid' arrows indicate 'stochastic' relationships – i.e. subject to random variability;
- 'dotted' arrows indicate 'deterministic' relationships – i.e. a given input will always provide the same output;

Estimation (Point Estimate)

Define a loss function, and minimise!

- Likelihood: $\mathcal{L}_i(\mathbf{y} \mid \lambda)$
- for an observation i : $\mathcal{L}_i = \frac{e^{-\lambda} \lambda^{y_i}}{y_i!}$
- for the entire sample: $\mathcal{L} = \prod_i^n \mathcal{L}_i$
- log-likelihood: $\log(\mathcal{L}) = \log(\prod_i^n \mathcal{L}_i) = \sum_i^n \log(\mathcal{L}_i) \leftarrow$ this we want to maximise
- negative log-likelihood:
$$L = -\sum_i^n \log(\mathcal{L}_i) = -\sum_i^n (y_i \log(\lambda) - \lambda - \log(y_i!))$$

⇒ This has an analytical solution:

$$\hat{\lambda}_{MLE} = \frac{1}{n} \sum_i^n y_i = \bar{y} \leftarrow \text{the sample mean is the MLE of } \lambda.$$

Uncertainty

- by asymptotic normality ($n \rightarrow \infty$) / Central Limit Theorem / Laplace approximation:

$$\lambda \sim N\left(\hat{\lambda}_{MLE}, \frac{\hat{\lambda}_{MLE}}{n}\right)$$


- This is a good approximation of the true posterior distribution when:
 - ❶ n is large enough;
 - ❷ the true rate is substantially above 0.
- When the approximation fails, we risk producing samples estimates below 0...

- Alternative distribution by Bayesian empirical posterior distribution:

$$\lambda \sim \text{Gamma}(\alpha = \sum_i^n y_i, \beta = n);$$


- α : 'shape' – n. of observed events;
- β : 'rate' – n. of intervals under consideration ;

Gamma Distribution



Simulation-Based (Monte-Carlo) Inference

- What statements can be made about λ ?
- draw S values from the empirical posterior:
 $\lambda_s \sim N\left(\hat{\lambda}_{MLE}, \frac{\hat{\lambda}_{MLE}}{n}\right);$
 $\lambda_s \sim \text{Gamma}(\alpha = \sum_i^n y_i, \beta = n);$
 $\lambda_{1:S} = \{0.24, 0.15, 0.12, 0.03, 0.04 \dots\}^3;$
- Plot the distribution of λ and infer its properties:



³These are all below 0 because for this example I chose a very rare event rate to display when the normal approximation can go wrong, but in principle they need not be !

Prediction


Posterior Predictive Distribution:

- 1 simulate S 'new' values from λ according to its empirical posterior distribution:

$$\tilde{\lambda}_s \sim \text{Gamma}(\alpha = \sum_i^n y_i, \beta = n)$$

- 2 simulate S new values \tilde{y} according to its likelihood, conditional on the simulated values of λ :

$$\tilde{y}_s \sim \text{Poisson}(\tilde{\lambda}_s)$$



The 'Heterogeneous Rate' Poisson Model

$$y_i \sim \text{Poisson}(\lambda_i)$$

$$\forall y_i \in \mathbb{N}_0, i \in \{1, \dots, n\},$$


$$\lambda_i = e^{\mu_i}$$

$$\mu_i = \beta_0 + \sum_{j=1}^p \beta_j x_{ij}$$

$$\lambda_i = e^{\mu_i} \rightarrow \text{log link function}$$

- ⇒ μ_i is transformed to a rate λ :
- ☎ Regression coefficients are interpretable in the same way as log-linear regression !
- ⇒ λ is itself interpretable as the expected rate of at which an event occurs;
- ✓ the Poisson distribution can natively handle $y = 0$, so no need to transform the data...

The 'Heterogeneous Rate' Poisson Model



Estimation (Point Estimate)

Define a loss function, and minimise !

- Likelihood for Poisson regression: $\mathcal{L}_i(y_i \mid \beta_0, \dots, \beta_p)$
- For an observation i : $\mathcal{L}_i = \frac{\exp(\beta_0 + \sum_{j=1}^p \beta_j x_{ij})^{y_i} \exp(-\exp(\beta_0 + \sum_{j=1}^p \beta_j x_{ij}))}{y_i!}$
- Minimise:

$$L = - \sum_{i=1}^n \log(\mathcal{L}_i) = - \sum_{i=1}^n \left[y_i(\beta_0 + \sum_{j=1}^p \beta_j x_{ij}) - \exp(\beta_0 + \sum_{j=1}^p \beta_j x_{ij}) - \log(y_i!) \right]$$

⇒ This **does not have an analytical solution** !

- by asymptotic normality ($n \rightarrow \infty$) / Laplace approximation:

$$\beta \sim MN(\hat{\beta}, \hat{\Sigma})$$

- software will estimate $\hat{\beta}$ via optimisation...
- use estimate $\hat{\Sigma}$ by plugging in $\hat{\beta}$ into the *Hessian* matrix...

Posterior Predictive Distribution

- ❶ define a set of L 'new subjects' characterised by design vector:
 $\tilde{x}_l, \forall l \in \{1, \dots, L\};$
- ❷ simulate S 'new' values from β according to its empirical (joint) posterior distribution:
 $\beta^s \sim N(\hat{\beta}_{MLE}, \hat{\Sigma});$
- ❸ Calculate μ_l (log-scale rate) for each simulation round:
 $\tilde{\mu}_l^s = \beta_0^s + \beta_1^s \tilde{x}_{l1} + \dots + \beta_p^s \tilde{x}_{lp}$
- ❹ Calculate λ_l (rate / expected count) for each simulation round:
 $\tilde{\lambda}_l^s = \exp(\mu_l^s)$
- ❺ simulate S new values \tilde{y} according to its likelihood, conditional on the simulated values of $\tilde{\lambda}$:
 $\tilde{y}_l^s \sim \text{Poisson}(\tilde{\lambda}_l^s);$
 - You can then use Monte Carlo methods to make inference about predictions – the MC median will typically be 'your best guess' for the expected count.